In this assignment, you will implement a matrix-matrix multiplication kernel that uses shared memory tiling. Your kernel is expected to work for any set of matrix dimensions so make sure to handle boundary conditions correctly.

Instructions

1. Place the files provided with this assignment in a single directory. The files are:
   - `main.cu`: contains setup and sequential code
   - `kernel.cu`: where you will implement your code (you should only modify this file)
   - `common.h`: for shared declarations across `main.cu` and `kernel.cu`
   - `timer.h`: to assist with timing
   - `Makefile`: used for compilation

2. Edit `kernel.cu` where TODO is indicated to implement the following:
   - Allocate device memory
   - Copy data from the host to the device
   - Configure and invoke the CUDA kernel
   - Copy the results from the device to the host
   - Free device memory
   - Perform the computation in the kernel

3. Compile your code by running: `make`

4. Test your code by running: `./mm-tiled`
   - If you are using the HPC cluster, do not forget to use the submission system. Do not run on the head node!
   - For testing on different matrix sizes, you can provide your own values for matrix dimensions as follows: `./mm-tiled <M> <N> <K>`

5. You are also provided with a file called `questions.txt` which contains questions about the assignment. Answer the questions in the file after implementing your kernel.

Submission
Submit your modified `kernel.cu` and `questions.txt` files via Moodle by the due date. Do not submit any other files or compressed folders.